

```

1  import os # Importing the os module to interact with the operating system
2  import numpy as np # Importing NumPy for numerical operations
3  import tensorflow as tf # Importing TensorFlow for deep learning
4  from tensorflow.keras import layers, models # Importing Keras layers and
models
5  from sklearn.model_selection import train_test_split # Importing function to
split data into training and validation sets
6  from sklearn.metrics import f1_score # Importing F1 score metric for
evaluation
7  import cv2 # Importing OpenCV for image processing
8  from tensorflow.keras.applications import MobileNetV2 # Importing MobileNetV2
model for transfer learning
9  from tensorflow.keras.layers import Dense, GlobalAveragePooling2D # Importing
Dense and Global Average Pooling layers
10 from tensorflow.keras.preprocessing.image import ImageDataGenerator #
Importing class for image data augmentation
11 from sklearn.utils import class_weight # Importing function to compute class
weights
12 from tensorflow.keras.callbacks import EarlyStopping, ReduceLROnPlateau,
ModelCheckpoint # Importing callbacks for training
13
14 # Check and configure GPU usage
15 if tf.config.list_physical_devices('GPU'): # Check if a GPU is available
16     print("GPU is available. Using CUDA for training.") # If GPU is available,
notify that CUDA will be used
17 else:
18     print("GPU is not available. Training will be performed on the CPU.") # If
not, notify that CPU will be used
19
20
21 # Dataset loading and Preprocessing Phase
22 # Constants
23 img_dir = '/content/drive/My Drive/processed_images/' # Directory where
processed images are stored
24 lbl_dir = '/content/drive/My Drive/processed_labels/' # Directory where
corresponding labels are stored
25 TARGET_CLASSES = ['Car', 'Cyclist', 'Pedestrian', 'Van', 'Truck'] # List of
target classes for classification
26 image_size = (224, 224) # Target image size for model input
27
28 # Load images and labels
29 def load_data(img_dir, lbl_dir): # Define function to load images and labels
30     images, labels = [], [] # Initialize empty lists for images and labels
31     for image_file in sorted(os.listdir(img_dir)): # Loop through sorted image
files in the image directory
32         img = cv2.imread(os.path.join(img_dir, image_file)) # Read the image
file
33         img = cv2.resize(img, image_size) / 255.0 # Resize the image and
normalize pixel values to [0, 1]

```

```

34         images.append(img) # Append the processed image to the images list
35
36         label_file = os.path.splitext(image_file)[0] + '.txt' # Construct the
label file name based on the image file name
37         label_path = os.path.join(lbl_dir, label_file) # Create the full path
for the label file
38         if os.path.exists(label_path): # Check if the label file exists
39             with open(label_path, 'r') as f: # Open the label file
40                 for line in f: # Loop through each line in the label file
41                     class_label = line.strip().split()[0] # Extract the class
label from the line
42                     if class_label in TARGET_CLASSES: # Check if the class
label is one of the target classes
43                         labels.append(TARGET_CLASSES.index(class_label)) #
Append the corresponding index of the class label to labels
44                     else:
45                         print(f"Warning: Missing label file for {image_file}") # Warn if
the label file is missing
46
47         return np.array(images), np.array(labels) # Return the images and labels
as NumPy arrays
48
49 # Load data
50 X, y = load_data(img_dir, lbl_dir) # Call the load_data function to load
images and labels
51 print(f"Initial shapes → X: {X.shape}, y: {y.shape}") # Print the shapes of
the loaded images and labels
52
53 # Ensure that y is a valid length corresponding to X
54 if len(X) < len(y): # Check if the number of images is less than the number of
labels
55     print("Warning: The length of y exceeds the length of X. Adjusting y to
match X.") # Warn if y exceeds X
56     y = y[:len(X)] # Trim y to match the length of X
57
58 # Filter out invalid labels
59 valid_indices = [i for i in range(len(X)) if y[i] is not None] # Create a list
of valid indices where y is not None
60 X_filtered = X[valid_indices] # Filter images based on valid indices
61 y_filtered = y[valid_indices] # Filter labels based on valid indices
62
63 # Convert labels to categorical
64 y_filtered = tf.keras.utils.to_categorical(y_filtered,
num_classes=len(TARGET_CLASSES)) # Convert labels to one-hot encoded format
65
66 print(f"Filtered shapes → X_filtered: {X_filtered.shape}, y_filtered:
{y_filtered.shape}") # Print the shapes of the filtered images and labels
67
68 # Split data into training and validation sets
69 X_train, X_val, y_train, y_val = train_test_split(X_filtered, y_filtered,
test_size=0.2, random_state=42) # Split data into training and validation sets

```

with a 80-20 ratio

```
70
71 # Augment Data
72 datagen = ImageDataGenerator( # Initialize ImageDataGenerator for data
    augmentation
73     rotation_range=20, # Randomly rotate images by up to 20 degrees
74     width_shift_range=0.2, # Randomly shift images horizontally by up to 20%
75     height_shift_range=0.2, # Randomly shift images vertically by up to 20%
76     horizontal_flip=True, # Randomly flip images horizontally
77     zoom_range=0.2, # Randomly zoom images by up to 20%
78     fill_mode='nearest' # Fill in new pixels after transformations
79 )
80
81
82
83 # Model define and Training
84 # Define MobileNetV2 model
85 base_model = MobileNetV2(weights='imagenet', include_top=False, input_shape=
    (224, 224, 3)) # Load MobileNetV2 without the top layer for transfer learning
86
87 for layer in base_model.layers[:-20]: # Freeze most layers in the base model
    for transfer learning
88     layer.trainable = False # Set the trainable property to False to prevent
    weight updates during training
89
90 # Add classifier layers
91 model = tf.keras.Sequential([ # Initialize a Sequential model
92     base_model, # Add the base model (MobileNetV2)
93     GlobalAveragePooling2D(), # Add global average pooling to reduce
    dimensionality
94     Dense(256, activation='relu'), # Add a fully connected layer with ReLU
    activation
95     Dense(len(TARGET_CLASSES), activation='softmax') # Add output layer with
    softmax activation for multi-class classification
96 ])
97
98 # Compile model
99 optimizer = tf.keras.optimizers.Adam(learning_rate=1e-5) # Initialize Adam
    optimizer with a low learning rate
100 model.compile(optimizer=optimizer, loss='categorical_crossentropy', metrics=
    ['accuracy']) # Compile the model with the optimizer, loss function, and
    metrics
101
102 # Class weights
103 class_weights = class_weight.compute_class_weight('balanced',
    classes=np.unique(np.argmax(y_train, axis=1)), y=np.argmax(y_train, axis=1)) #
    Compute class weights to handle class imbalance
104 class_weights_dict = dict(enumerate(class_weights)) # Convert class weights to
    a dictionary format for use in training
105
106 # Callbacks
```

```

107 early_stopping = EarlyStopping(monitor='val_loss', patience=10,
108 restore_best_weights=True) # Early stopping to prevent overfitting
109 reduce_lr = ReduceLROnPlateau(monitor='val_loss', factor=0.5, patience=5,
110 min_lr=1e-6) # Reduce learning rate when validation loss plateaus
111 model_checkpoint = ModelCheckpoint('best_model.keras', save_best_only=True,
112 monitor='val_loss') # Save the best model based on validation loss
113
114 # Training setup
115 batch_size = 8 # Set the batch size for training
116 train_gen = datagen.flow(X_train, y_train, batch_size=batch_size) # Create a
117 generator for training data augmentation
118
119 # Train the model
120 history = model.fit( # Fit the model to the training data
121     train_gen, # Use the augmented training generator
122     steps_per_epoch=len(X_train) // batch_size, # Calculate steps per epoch
123     based on batch size
124     epochs=100, # Set the number of training epochs
125     validation_data=(X_val, y_val), # Validate the model on the validation set
126     class_weight=class_weights_dict, # Use computed class weights during
127     training
128     callbacks=[early_stopping, reduce_lr, model_checkpoint] # Include
129     callbacks for training
130 )
131
132 # Evaluate the model
133 y_pred = np.argmax(model.predict(X_val), axis=1) # Use the model to predict
134 classes on the validation set and get the class indices with the highest
135 probabilities
136 y_true = np.argmax(y_val, axis=1) # Convert one-hot encoded true labels back
137 to class indices
138
139 # Compute F1 scores for all classes
140 f1_per_class = f1_score(y_true, y_pred, average=None,
141 labels=np.arange(len(TARGET_CLASSES))) # Compute F1 scores for each class
142 separately, specifying the target class labels
143 f1_weighted = f1_score(y_true, y_pred, average='weighted') # Compute the
144 weighted F1 score, considering the support of each class
145
146 # Print F1 scores
147 for i, class_name in enumerate(TARGET_CLASSES): # Loop through each target
148 class
149     if i < len(f1_per_class): # Check if F1 score is available for the class
150         print(f"F1 Score for {class_name}: {f1_per_class[i]:.4f}") # Print the
151 F1 score for the class formatted to four decimal places
152     else:
153         print(f"F1 Score for {class_name}: Not applicable (no instances
154 predicted)") # Notify if there are no predicted instances for the class

```

```

142 print("Weighted F1 Score:", f1_weighted) # Print the overall weighted F1 score
143
144
145
146
147 #Model testing phase.
148 import os
149 import numpy as np
150 import cv2
151 import tensorflow as tf
152 from tensorflow.keras.applications import MobileNetV2
153 from tensorflow.keras.layers import Dense, GlobalAveragePooling2D
154 from tensorflow.keras.models import Model
155 import matplotlib.pyplot as plt
156
157 # Constants
158 TARGET_CLASSES = ['Car', 'Cyclist', 'Pedestrian', 'Van', 'Truck'] # List of
target classes for classification
159 image_size = (224, 224) # Target size for images
160
161 # Function to build the model
162 def build_model():
163     # Load the MobileNetV2 model without the top layer, using pretrained
ImageNet weights
164     base_model = MobileNetV2(weights='imagenet', include_top=False,
input_shape=(224, 224, 3))
165     for layer in base_model.layers[:-20]: # Freeze most layers to retain pre-
trained weights
166         layer.trainable = False
167
168     # Add a global average pooling layer and a dense layer for classification
169     x = GlobalAveragePooling2D()(base_model.output)
170     x = Dense(256, activation='relu')(x) # Fully connected layer with ReLU
activation
171     outputs = Dense(len(TARGET_CLASSES), activation='softmax')(x) # Output
layer for multi-class classification
172
173     model = Model(inputs=base_model.input, outputs=outputs) # Create the model
with specified inputs and outputs
174     return model
175
176 # Load the model architecture
177 model = build_model() # Instantiate the model
178
179 # Load the saved weights
180 try:
181     model.load_weights('best_model.keras') # Load the weights from the saved
file
182     print("Weights loaded successfully.") # Confirmation message
183 except Exception as e:
184     print("Error loading weights:", e) # Error message if loading fails

```

```

185
186 # Print model summary to check its structure
187 model.summary() # Display the architecture of the model
188
189 # Function to load and preprocess test images
190 def load_test_data(test_dir):
191     images = [] # Initialize list for images
192     image_files = [] # Initialize list for image filenames
193     for image_file in sorted(os.listdir(test_dir)): # Iterate over files in
the test directory
194         img = cv2.imread(os.path.join(test_dir, image_file)) # Read the image
file
195         if img is None: # Check if the image was loaded successfully
196             print(f"Warning: {image_file} could not be read. Skipping.") #
Warning for unreadable images
197             continue # Skip to the next file
198         img = cv2.resize(img, image_size) / 255.0 # Resize and normalize the
image
199         images.append(img) # Add image to list
200         image_files.append(image_file) # Store image filename for later
201     return np.array(images), image_files # Return numpy array of images and
list of filenames
202
203 # Load test data
204 test_dir = '/content/drive/My Drive/test_images/' # Directory containing test
images
205 X_test, image_files = load_test_data(test_dir) # Load images and filenames
206
207 # Ensure the shape is correct
208 print(f"Test data shape: {X_test.shape}") # Print the shape of the loaded test
data
209
210 # Define bounding box scaling factors for each class
211 bounding_box_scale_factors = {
212     'Car': (0.4, 0.3), # Scale factors for bounding box dimensions for each
class
213     'Cyclist': (0.2, 0.4),
214     'Pedestrian': (0.1, 0.5),
215     'Van': (0.5, 0.3),
216     'Truck': (0.5, 0.4),
217 }
218
219 # Run predictions only if there are images to predict
220 if X_test.shape[0] > 0: # Check if there are any test images
221     y_pred_probs = model.predict(X_test) # Get prediction probabilities for
the test images
222     y_pred = np.argmax(y_pred_probs, axis=1) # Determine the predicted class
indices
223
224     # Map predictions to class labels

```

```

225     predicted_labels = [TARGET_CLASSES[i] for i in y_pred] # Convert indices
to class names
226
227     # Display predicted labels for each test image
228     for filename, label in zip(image_files, predicted_labels): # Iterate
through filenames and predicted labels
229         print(f"{filename}: {label}") # Print the filename and its predicted
label
230
231         # Load original image to draw the bounding box
232         img = cv2.imread(os.path.join(test_dir, filename)) # Load the original
image
233         height, width, _ = img.shape # Get image dimensions
234
235         # Get bounding box scale based on the predicted class
236         scale_x, scale_y = bounding_box_scale_factors[label] # Retrieve scale
factors for the predicted class
237
238         # Calculate bounding box coordinates
239         box_width = int(width * scale_x) # Calculate the bounding box width
240         box_height = int(height * scale_y) # Calculate the bounding box height
241
242         # Center the bounding box in the image
243         x_min = (width - box_width) // 2 # Calculate the minimum x-coordinate
244         y_min = (height - box_height) // 2 # Calculate the minimum y-
coordinate
245         x_max = x_min + box_width # Calculate the maximum x-coordinate
246         y_max = y_min + box_height # Calculate the maximum y-coordinate
247
248         # Draw bounding box and label on the image
249         cv2.rectangle(img, (x_min, y_min), (x_max, y_max), (0, 0, 255), 2) #
Draw a red rectangle for the bounding box
250         cv2.putText(img, label, (x_min, y_min - 10), cv2.FONT_HERSHEY_SIMPLEX,
0.5, (0, 0, 255), 2) # Add the label text above the bounding box
251
252         # Display the image with bounding box
253         plt.imshow(cv2.cvtColor(img, cv2.COLOR_BGR2RGB)) # Convert BGR to RGB
for displaying
254         plt.axis('off') # Turn off the axis
255         plt.title(f"Predicted: {label}") # Set the title to show the predicted
label
256         plt.show() # Display the image
257     else:
258         print("No valid test images to predict.") # Message if no valid test
images were found
259
260

```